

tags: master, kube-apiserver

05-2. 部署 kube-apiserver 集群

- 05-2. 部署 kube-apiserver 集群
 - 创建 kubernetes-master 证书和私钥
 - 创建加密配置文件
 - 创建审计策略文件
 - 创建后续访问 metrics-server 或 kube-prometheus 使用的证书
 - 创建 kube-apiserver systemd unit 模板文件
 - 为各节点创建和分发 kube-apiserver systemd unit 文件
 - 启动 kube-apiserver 服务
 - 检查 kube-apiserver 运行状态
 - 检查集群信息
 - 检查 kube-apiserver 监听的端口

本文档讲解部署一个三实例 kube-apiserver 集群的步骤。

注意：如果没有特殊指明，本文档的所有操作均在 **zhangjun-k8s-01** 节点上执行。

创建 kubernetes-master 证书和私钥

创建证书签名请求：

```
cd /opt/k8s/work
source /opt/k8s/bin/environment.sh
cat > kubernetes-csr.json <<EOF
{
  "CN": "kubernetes-master",
  "hosts": [
    "127.0.0.1",
    "172.27.138.251",
    "172.27.137.229",
    "172.27.138.239",
    "${CLUSTER_KUBERNETES_SVC_IP}",
    "kubernetes",
    "kubernetes.default",
    "kubernetes.default.svc",
    "kubernetes.default.svc.cluster",
    "kubernetes.default.svc.cluster.local.",
    "kubernetes.default.svc.${CLUSTER_DNS_DOMAIN}."
  ],
  "isCA": false,
  "keyUsage": [
    "digitalSignature",
    "keyEncipherment"
  ],
  "signatureAlgorithm": "sha256WithRSAEncryption",
  "validity": {
    "notAfter": "2025-12-31T23:59:59Z",
    "notBefore": "2025-01-01T00:00:00Z"
  }
}
```

```

"key": {
  "algo": "rsa",
  "size": 2048
},
"names": [
  {
    "C": "CN",
    "ST": "BeiJing",
    "L": "BeiJing",
    "O": "k8s",
    "OU": "opsnull"
  }
]
}
EOF

```

- hosts 字段指定授权使用该证书的 IP 和域名列表，这里列出了 master 节点 IP、kubernetes 服务的 IP 和域名；

生成证书和私钥：

```

cfssl gencert -ca=/opt/k8s/work/ca.pem \
  -ca-key=/opt/k8s/work/ca-key.pem \
  -config=/opt/k8s/work/ca-config.json \
  -profile=kubernetes kubernetes-csr.json | cfssljson -bare kubernetes
ls kubernetes*.pem

```

将生成的证书和私钥文件拷贝到所有 master 节点：

```

cd /opt/k8s/work
source /opt/k8s/bin/environment.sh
for node_ip in ${NODE_IPS[@]}
do
  echo ">>> ${node_ip}"
  ssh root@${node_ip} "mkdir -p /etc/kubernetes/cert"
  scp kubernetes*.pem root@${node_ip}:/etc/kubernetes/cert/
done

```

创建加密配置文件

```

cd /opt/k8s/work
source /opt/k8s/bin/environment.sh
cat > encryption-config.yaml <<EOF
kind: EncryptionConfig
apiVersion: v1
resources:
  - resources:
    - secrets

```

```
providers:
  - aescbc:
      keys:
        - name: key1
          secret: ${ENCRYPTION_KEY}
  - identity: {}
EOF
```

将加密配置文件拷贝到 master 节点的 /etc/kubernetes 目录下:

```
cd /opt/k8s/work
source /opt/k8s/bin/environment.sh
for node_ip in ${NODE_IPS[@]}
do
  echo ">>> ${node_ip}"
  scp encryption-config.yaml root@${node_ip}:/etc/kubernetes/
done
```

创建审计策略文件

```
cd /opt/k8s/work
source /opt/k8s/bin/environment.sh
cat > audit-policy.yaml <<EOF
apiVersion: audit.k8s.io/v1beta1
kind: Policy
rules:
  # The following requests were manually identified as high-volume and low-risk, so
  drop them.
  - level: None
    resources:
      - group: ""
        resources:
          - endpoints
          - services
          - services/status
    users:
      - 'system:kube-proxy'
    verbs:
      - watch

  - level: None
    resources:
      - group: ""
        resources:
          - nodes
          - nodes/status
    userGroups:
      - 'system:nodes'
```

```
verbs:
  - get

- level: None
  namespaces:
    - kube-system
  resources:
    - group: ""
      resources:
        - endpoints
  users:
    - 'system:kube-controller-manager'
    - 'system:kube-scheduler'
    - 'system:serviceaccount:kube-system:endpoint-controller'
  verbs:
    - get
    - update

- level: None
  resources:
    - group: ""
      resources:
        - namespaces
        - namespaces/status
        - namespaces/finalize
  users:
    - 'system:apiserver'
  verbs:
    - get

# Don't log HPA fetching metrics.
- level: None
  resources:
    - group: metrics.k8s.io
  users:
    - 'system:kube-controller-manager'
  verbs:
    - get
    - list

# Don't log these read-only URLs.
- level: None
  nonResourceURLs:
    - '/healthz*'
    - '/version'
    - '/swagger*'

# Don't log events requests.
- level: None
  resources:
    - group: ""
      resources:
```

- events

node and pod status calls from nodes are high-volume and can be large, don't log responses

for expected updates from nodes

- level: Request
 - omitStages:
 - RequestReceived
 - resources:
 - group: ""
 - resources:
 - nodes/status
 - pods/status
 - users:
 - kubelet
 - 'system:node-problem-detector'
 - 'system:serviceaccount:kube-system:node-problem-detector'
 - verbs:
 - update
 - patch

- level: Request
 - omitStages:
 - RequestReceived
 - resources:
 - group: ""
 - resources:
 - nodes/status
 - pods/status
 - userGroups:
 - 'system:nodes'
 - verbs:
 - update
 - patch

deletecollection calls can be large, don't log responses for expected namespace deletions

- level: Request
 - omitStages:
 - RequestReceived
 - users:
 - 'system:serviceaccount:kube-system:namespace-controller'
 - verbs:
 - deletecollection

Secrets, ConfigMaps, and TokenReviews can contain sensitive & binary data, # so only log at the Metadata level.

- level: Metadata
 - omitStages:
 - RequestReceived
 - resources:
 - group: ""

```

    resources:
      - secrets
      - configmaps
    - group: authentication.k8s.io
      resources:
        - tokenreviews
# Get repsonses can be large; skip them.
- level: Request
  omitStages:
    - RequestReceived
  resources:
    - group: ""
    - group: admissionregistration.k8s.io
    - group: apiextensions.k8s.io
    - group: apiregistration.k8s.io
    - group: apps
    - group: authentication.k8s.io
    - group: authorization.k8s.io
    - group: autoscaling
    - group: batch
    - group: certificates.k8s.io
    - group: extensions
    - group: metrics.k8s.io
    - group: networking.k8s.io
    - group: policy
    - group: rbac.authorization.k8s.io
    - group: scheduling.k8s.io
    - group: settings.k8s.io
    - group: storage.k8s.io
  verbs:
    - get
    - list
    - watch

# Default level for known APIs
- level: RequestResponse
  omitStages:
    - RequestReceived
  resources:
    - group: ""
    - group: admissionregistration.k8s.io
    - group: apiextensions.k8s.io
    - group: apiregistration.k8s.io
    - group: apps
    - group: authentication.k8s.io
    - group: authorization.k8s.io
    - group: autoscaling
    - group: batch
    - group: certificates.k8s.io
    - group: extensions
    - group: metrics.k8s.io
    - group: networking.k8s.io

```

```

- group: policy
- group: rbac.authorization.k8s.io
- group: scheduling.k8s.io
- group: settings.k8s.io
- group: storage.k8s.io

# Default level for all other requests.
- level: Metadata
  omitStages:
    - RequestReceived
EOF

```

分发审计策略文件:

```

cd /opt/k8s/work
source /opt/k8s/bin/environment.sh
for node_ip in ${NODE_IPS[@]}
do
    echo ">>> ${node_ip}"
    scp audit-policy.yaml root@${node_ip}:/etc/kubernetes/audit-policy.yaml
done

```

创建后续访问 metrics-server 或 kube-prometheus 使用的证书

创建证书签名请求:

```

cd /opt/k8s/work
cat > proxy-client-csr.json <<EOF
{
  "CN": "aggregator",
  "hosts": [],
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "CN",
      "ST": "BeiJing",
      "L": "BeiJing",
      "O": "k8s",
      "OU": "opsnull"
    }
  ]
}
EOF

```

- CN 名称需要位于 kube-apiserver 的 `--requestheader-allowed-names` 参数中, 否则后续访问 metrics 时会

提示权限不足。

生成证书和私钥：

```
cfssl gencert -ca=/etc/kubernetes/cert/ca.pem \  
-ca-key=/etc/kubernetes/cert/ca-key.pem \  
-config=/etc/kubernetes/cert/ca-config.json \  
-profile=kubernetes proxy-client-csr.json | cfssljson -bare proxy-client  
ls proxy-client*.pem
```

将生成的证书和私钥文件拷贝到所有 master 节点：

```
source /opt/k8s/bin/environment.sh  
for node_ip in ${NODE_IPS[@]}  
do  
    echo ">>> ${node_ip}"  
    scp proxy-client*.pem root@${node_ip}:/etc/kubernetes/cert/  
done
```

创建 kube-apiserver systemd unit 模板文件

```
cd /opt/k8s/work  
source /opt/k8s/bin/environment.sh  
cat > kube-apiserver.service.template <<EOF  
[Unit]  
Description=Kubernetes API Server  
Documentation=https://github.com/GoogleCloudPlatform/kubernetes  
After=network.target  
  
[Service]  
WorkingDirectory=${K8S_DIR}/kube-apiserver  
ExecStart=/opt/k8s/bin/kube-apiserver \  
--advertise-address=##NODE_IP## \  
--default-not-ready-toleration-seconds=360 \  
--default-unreachable-toleration-seconds=360 \  
--feature-gates=DynamicAuditing=true \  
--max-mutating-requests-inflight=2000 \  
--max-requests-inflight=4000 \  
--default-watch-cache-size=200 \  
--delete-collection-workers=2 \  
--encryption-provider-config=/etc/kubernetes/encryption-config.yaml \  
--etcd-cafile=/etc/kubernetes/cert/ca.pem \  
--etcd-certfile=/etc/kubernetes/cert/kubernetes.pem \  
--etcd-keyfile=/etc/kubernetes/cert/kubernetes-key.pem \  
--etcd-servers=${ETCD_ENDPOINTS} \  
--bind-address=##NODE_IP## \  
--secure-port=6443 \  
--tls-cert-file=/etc/kubernetes/cert/kubernetes.pem \  

```



```

--tls-private-key-file=/etc/kubernetes/cert/kubernetes-key.pem \\  

--insecure-port=0 \\  

--audit-dynamic-configuration \\  

--audit-log-maxage=15 \\  

--audit-log-maxbackup=3 \\  

--audit-log-maxsize=100 \\  

--audit-log-truncate-enabled \\  

--audit-log-path=${K8S_DIR}/kube-apiserver/audit.log \\  

--audit-policy-file=/etc/kubernetes/audit-policy.yaml \\  

--profiling \\  

--anonymous-auth=false \\  

--client-ca-file=/etc/kubernetes/cert/ca.pem \\  

--enable-bootstrap-token-auth \\  

--requestheader-allowed-names="aggregator" \\  

--requestheader-client-ca-file=/etc/kubernetes/cert/ca.pem \\  

--requestheader-extra-headers-prefix="X-Remote-Extra-" \\  

--requestheader-group-headers=X-Remote-Group \\  

--requestheader-username-headers=X-Remote-User \\  

--service-account-key-file=/etc/kubernetes/cert/ca.pem \\  

--authorization-mode=Node,RBAC \\  

--runtime-config=api/all=true \\  

--enable-admission-plugins=NodeRestriction \\  

--allow-privileged=true \\  

--apiserver-count=3 \\  

--event-ttl=168h \\  

--kubelet-certificate-authority=/etc/kubernetes/cert/ca.pem \\  

--kubelet-client-certificate=/etc/kubernetes/cert/kubernetes.pem \\  

--kubelet-client-key=/etc/kubernetes/cert/kubernetes-key.pem \\  

--kubelet-https=true \\  

--kubelet-timeout=10s \\  

--proxy-client-cert-file=/etc/kubernetes/cert/proxy-client.pem \\  

--proxy-client-key-file=/etc/kubernetes/cert/proxy-client-key.pem \\  

--service-cluster-ip-range=${SERVICE_CIDR} \\  

--service-node-port-range=${NODE_PORT_RANGE} \\  

--logtostderr=true \\  

--v=2  

Restart=on-failure  

RestartSec=10  

Type=notify  

LimitNOFILE=65536  
  

[Install]  

WantedBy=multi-user.target  

EOF

```

- --advertise-address: apiserver 对外通告的 IP (kubernetes 服务后端节点 IP) ;
- --default-*--toleration-seconds: 设置节点异常相关的阈值;
- --max-*--requests-inflight: 请求相关的最大阈值;
- --etcd-*: 访问 etcd 的证书和 etcd 服务器地址;
- --bind-address: https 监听的 IP, 不能为 127.0.0.1, 否则外界不能访问它的安全端口 6443;
- --secret-port: https 监听端口;

- `--insecure-port=0`: 关闭监听 http 非安全端口(8080);
- `--tls-*--file`: 指定 apiserver 使用的证书、私钥和 CA 文件;
- `--audit-*`: 配置审计策略和审计日志文件相关的参数;
- `--client-ca-file`: 验证 client (kue-controller-manager、kube-scheduler、kubelet、kube-proxy 等)请求所带的证书;
- `--enable-bootstrap-token-auth`: 启用 kubelet bootstrap 的 token 认证;
- `--requestheader-*`: kube-apiserver 的 aggregator layer 相关的配置参数, proxy-client & HPA 需要使用;
- `--requestheader-client-ca-file`: 用于签名 `--proxy-client-cert-file` 和 `--proxy-client-key-file` 指定的证书; 在启用了 metric aggregator 时使用;
- `--requestheader-allowed-names`: 不能为空, 值为逗号分割的 `--proxy-client-cert-file` 证书的 CN 名称, 这里设置为 "aggregator";
- `--service-account-key-file`: 签名 ServiceAccount Token 的公钥文件, kube-controller-manager 的 `--service-account-private-key-file` 指定私钥文件, 两者配对使用;
- `--runtime-config=api/all=true`: 启用所有版本的 APIs, 如 autoscaling/v2alpha1;
- `--authorization-mode=Node,RBAC`、`--anonymous-auth=false`: 开启 Node 和 RBAC 授权模式, 拒绝未授权的请求;
- `--enable-admission-plugins`: 启用一些默认关闭的 plugins;
- `--allow-privileged`: 运行执行 privileged 权限的容器;
- `--apiserver-count=3`: 指定 apiserver 实例的数量;
- `--event-ttl`: 指定 events 的保存时间;
- `--kubelet-*`: 如果指定, 则使用 https 访问 kubelet APIs; 需要为证书对应的用户(上面 `kubernetes*.pem` 证书的用户为 `kubernetes`)用户定义 RBAC 规则, 否则访问 kubelet API 时提示未授权;
- `--proxy-client-*`: apiserver 访问 metrics-server 使用的证书;
- `--service-cluster-ip-range`: 指定 Service Cluster IP 地址段;
- `--service-node-port-range`: 指定 NodePort 的端口范围;

如果 kube-apiserver 机器没有运行 kube-proxy, 则还需要添加 `--enable-aggregator-routing=true` 参数;

关于 `--requestheader-XXX` 相关参数, 参考:

- <https://github.com/kubernetes-incubator/apiserver-builder/blob/master/docs/concepts/auth.md>
- <https://docs.bitnami.com/kubernetes/how-to/configure-autoscaling-custom-metrics/>

注意:

1. `--requestheader-client-ca-file` 指定的 CA 证书, 必须具有 client auth and server auth;
2. 如果 `--requestheader-allowed-names` 不为空, 且 `--proxy-client-cert-file` 证书的 CN 名称不在 `allowed-names` 中, 则后续查看 node 或 pods 的 metrics 失败, 提示:

```
$ kubectl top nodes
Error from server (Forbidden): nodes.metrics.k8s.io is forbidden: User "aggregator"
cannot list resource "nodes" in API group "metrics.k8s.io" at the cluster scope
```

为各节点创建和分发 kube-apiserver systemd unit 文件

替换模板文件中的变量, 为各节点生成 systemd unit 文件:

```
cd /opt/k8s/work
source /opt/k8s/bin/environment.sh
for (( i=0; i < 3; i++ ))
do
    sed -e "s/##NODE_NAME##/${NODE_NAMES[i]}/" -e "s/##NODE_IP##/${NODE_IPS[i]}/"
    kube-apiserver.service.template > kube-apiserver-${NODE_IPS[i]}.service
done
ls kube-apiserver*.service
```

- NODE_NAMES 和 NODE_IPS 为相同长度的 bash 数组，分别为节点名称和对应的 IP；

分发生成的 systemd unit 文件：

```
cd /opt/k8s/work
source /opt/k8s/bin/environment.sh
for node_ip in ${NODE_IPS[@]}
do
    echo ">>> ${node_ip}"
    scp kube-apiserver-${node_ip}.service root@${node_ip}:/etc/systemd/system/kube-
    apiserver.service
done
```

启动 kube-apiserver 服务

```
source /opt/k8s/bin/environment.sh
for node_ip in ${NODE_IPS[@]}
do
    echo ">>> ${node_ip}"
    ssh root@${node_ip} "mkdir -p ${K8S_DIR}/kube-apiserver"
    ssh root@${node_ip} "systemctl daemon-reload && systemctl enable kube-apiserver
    && systemctl restart kube-apiserver"
done
```

检查 kube-apiserver 运行状态

```
source /opt/k8s/bin/environment.sh
for node_ip in ${NODE_IPS[@]}
do
    echo ">>> ${node_ip}"
    ssh root@${node_ip} "systemctl status kube-apiserver |grep 'Active:'"
done
```

确保状态为 active (running)，否则查看日志，确认原因：

```
journalctl -u kube-apiserver
```

检查集群状态

```
$ kubectl cluster-info
```

Kubernetes master is running at https://172.27.138.251:6443

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

```
$ kubectl get all --all-namespaces
```

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default	service/kubernetes	ClusterIP	10.254.0.1	<none>	443/TCP	3m53s

```
$ kubectl get componentstatuses
```

NAME	AGE
controller-manager	<unknown>
scheduler	<unknown>
etcd-0	<unknown>
etcd-2	<unknown>
etcd-1	<unknown>

- Kubernetes 1.16.6 存在 Bugs 导致返回结果一直为 <unknown>, 但 `kubectl get cs -o yaml` 可以返回正确结果;

检查 kube-apiserver 监听的端口

```
$ sudo netstat -lnpt|grep kube
```

```
tcp        0      0 172.27.138.251:6443  0.0.0.0:*           LISTEN
101442/kube-apiserv
```

- 6443: 接收 https 请求的安全端口, 对所有请求做认证和授权;
- 由于关闭了非安全端口, 故没有监听 8080;