

tags: registry, ceph

09. 部署 docker registry

- 09. 部署 docker registry
 - 部署 ceph RGW 节点
 - 创建测试账号 demo
 - 创建 demo 账号的子账号 swift
 - 创建 demo:swift 子账号的 secret key
 - 创建 docker registry
 - 向 registry push image
 - 私有 registry 的运维操作
 - 查询私有镜像中的 images
 - 查询某个镜像的 tags 列表
 - 获取 image 或 layer 的 digest
 - 删除 image
 - 删除 layer
 - 常见问题
 - login 失败 416
 - login 失败 503

本文档介绍使用 docker 官方的 registry v2 镜像部署私有仓库的步骤，你也可以参考附件文档部署 Harbor 私有仓库（D.部署 Harbor 私有仓库）。

本文档讲解部署一个 TLS 加密、HTTP Basic 认证、用 ceph rgw 做后端存储的私有 docker registry 步骤，如果使用其它类型的后端存储，则可以从“创建 docker registry”节开始；

示例两台机器 IP 如下：

- ceph rgw: 172.27.132.66
- docker registry: 172.27.132.67

部署 ceph RGW 节点

```
$ ceph-deploy rgw create 172.27.132.66 # rgw 默认监听7480端口
$
```

创建测试账号 demo

```
$ radosgw-admin user create --uid=demo --display-name="ceph rgw demo user"
$
```

创建 demo 账号的子账号 swift

当前 registry 只支持使用 swift 协议访问 ceph rgw 存储，暂时不支持 s3 协议；

```
$ radosgw-admin subuser create --uid demo --subuser=demo:swift --access=full --secret=secretkey --key-type=swift
$
```

创建 demo:swift 子账号的 secret key

```
$ radosgw-admin key create --subuser=demo:swift --key-type=swift --gen-secret
{
  "user_id": "demo",
  "display_name": "ceph rgw demo user",
  "email": "",
  "suspended": 0,
  "max_buckets": 1000,
  "auid": 0,
  "subusers": [
    {
      "id": "demo:swift",
      "permissions": "full-control"
    }
  ],
  "keys": [
    {
      "user": "demo",
      "access_key": "5Y1B1SIJ2YHKEH05U36B",
      "secret_key": "nrIvtPqUj7pUlccLYPuR3ntVzIa50DToIpe7xFjT"
    }
  ],
  "swift_keys": [
    {
      "user": "demo:swift",
      "secret_key": "ttQcU1017DFQ4I9xzKqwgUe7WIYYX99zhcIfU9vb"
    }
  ],
  "caps": [],
  "op_mask": "read, write, delete",
}
```

```

    "default_placement": "",
    "placement_tags": [],
    "bucket_quota": {
        "enabled": false,
        "max_size_kb": -1,
        "max_objects": -1
    },
    "user_quota": {
        "enabled": false,
        "max_size_kb": -1,
        "max_objects": -1
    },
    "temp_url_keys": []
}

```

- ttQcU1017DFQ4I9xzKqwgUe7WIYYX99zhcIfU9vb 为子账号 demo:swift 的 secret key;

创建 docker registry

创建 registry 使用的 x509 证书

```

$ mkdir -p registry/{auth,certs}
$ cat > registry-csr.json <<EOF
{
    "CN": "registry",
    "hosts": [
        "127.0.0.1",
        "172.27.132.67"
    ],
    "key": {
        "algo": "rsa",
        "size": 2048
    },
    "names": [
        {
            "C": "CN",
            "ST": "BeiJing",
            "L": "BeiJing",
            "O": "k8s",
            "OU": "opsnull"
        }
    ]
}
EOF
$ cfssl gencert -ca=/etc/kubernetes/cert/ca.pem \
    -ca-key=/etc/kubernetes/cert/ca-key.pem \
    -config=/etc/kubernetes/cert/ca-config.json \
    -profile=kubernetes registry-csr.json | cfssljson -bare registry
$ cp registry.pem registry-key.pem registry/certs

```

```
$
```

- 这里复用以前创建的 CA 证书和秘钥文件;
- hosts 字段指定 registry 的 NodeIP;

创建 HTTP Basic 认证文件

```
$ docker run --entrypoint htpasswd registry:2 -Bbn foo foo123 >
registry/auth/htpasswd
$ cat registry/auth/htpasswd
foo:$2y$05$iZaM45Jxlcg0DJKXZMggL0ibAsHLGybyU.CgU9AHqWcVDyBjiScN.
```

配置 registry 参数

```
export RGW_AUTH_URL="http://172.27.132.66:7480/auth/v1"
export RGW_USER="demo:swift"
export RGW_SECRET_KEY="ttQcU1017DFQ4I9xzKqwgUe7WIYYX99zhcIfU9vb"
cat > config.yml << EOF
# https://docs.docker.com/registry/configuration/#list-of-configuration-options
version: 0.1
log:
  level: info
  formatter: text
  fields:
    service: registry

storage:
  cache:
    blobdescriptor: inmemory
  delete:
    enabled: true
  swift:
    authurl: ${RGW_AUTH_URL}
    username: ${RGW_USER}
    password: ${RGW_SECRET_KEY}
    container: registry

auth:
  htpasswd:
    realm: basic-realm
    path: /auth/htpasswd

http:
  addr: 0.0.0.0:8000
  headers:
    X-Content-Type-Options: [nosniff]
  tls:
    certificate: /certs/registry.pem
    key: /certs/registry-key.pem

health:
```

```
    storagedriver:
      enabled: true
      interval: 10s
      threshold: 3
EOF
[k8s@zhangjun-k8s-01 cert]$ cp config.yml registry
[k8s@zhangjun-k8s-01 cert]$ scp -r registry 172.27.132.67:/opt/k8s
```

- storage.swift 指定后端使用 swfit 接口协议的存储，这里配置的是 ceph rgw 存储参数；
- auth.htpasswd 指定了 HTTP Basic 认证的 token 文件路径；
- http.tls 指定了 registry http 服务器的证书和秘钥文件路径；

创建 docker registry：

```
ssh k8s@172.27.132.67
$ docker run -d -p 8000:8000 --privileged \
  -v /opt/k8s/registry/auth:/auth \
  -v /opt/k8s/registry/certs:/certs \
  -v /opt/k8s/registry/config.yml:/etc/docker/registry/config.yml \
  --name registry registry:2
```

- 执行该 docker run 命令的机器 IP 为 172.27.132.67；

向 registry push image

将签署 registry 证书的 CA 证书拷贝到 /etc/docker/certs.d/172.27.132.67:8000 目录下

```
[k8s@zhangjun-k8s-01 cert]$ sudo mkdir -p /etc/docker/certs.d/172.27.132.67:8000
[k8s@zhangjun-k8s-01 cert]$ sudo cp /etc/kubernetes/cert/ca.pem
/etc/docker/certs.d/172.27.132.67:8000/ca.crt
```

登陆私有 registry：

```
$ docker login 172.27.132.67:8000
Username: foo
Password:
Login Succeeded
```

登陆信息被写入 ~/.docker/config.json 文件：

```
$ cat ~/.docker/config.json
{
    "auths": {
        "172.27.132.67:8000": {
            "auth": "Zm9vOmZvbzEyMw=="
        }
    }
}
```

将本地的 image 打上私有 registry 的 tag:

```
$ docker tag prom/node-exporter:v0.16.0 172.27.132.67:8000/prom/node-exporter:v0.16.0
$ docker images |grep pause
prom/node-exporter:v0.16.0                                latest
f9d5de079539      2 years ago      239.8 kB
172.27.132.67:8000/prom/node-exporter:v0.16.0           latest
f9d5de079539      2 years ago      239.8 kB
```

将 image push 到私有 registry:

```
$ docker push 172.27.132.67:8000/prom/node-exporter:v0.16.0
The push refers to a repository [172.27.132.67:8000/prom/node-exporter:v0.16.0]
5f70bf18a086: Pushed
e16a89738269: Pushed
latest: digest:
sha256:9a6b437e896acad3f5a2a8084625fdd4177b2e7124ee943af642259f2f283359 size: 916
```

查看 ceph 上是否已经有 push 的 pause 容器文件:

```
$ rados lspools
rbd
cephfs_data
cephfs_metadata
.rgw.root
k8s
default.rgw.control
default.rgw.meta
default.rgw.log
default.rgw.buckets.index
default.rgw.buckets.data

$ rados --pool default.rgw.buckets.data ls|grep node-exporter
1f3f02c4-fe58-4626-992b-
c6c0fe4c8acf.34107.1_files/docker/registry/v2/repositories/prom/node-
exporter/_layers/sha256/cdb7590af5f064887f3d6008d46be65e929c74250d747813d85199e04fc70
463/link
```

```
1f3f02c4-fe58-4626-992b-  
c6c0fe4c8acf.34107.1_files/docker/registry/v2/repositories/prom/node-  
exporter/_manifests/revisions/sha256/55302581333c43d540db0e144cf9e7735423117a733cdec2  
7716d87254221086/link  
1f3f02c4-fe58-4626-992b-  
c6c0fe4c8acf.34107.1_files/docker/registry/v2/repositories/prom/node-  
exporter/_manifests/tags/v0.16.0/current/link  
1f3f02c4-fe58-4626-992b-  
c6c0fe4c8acf.34107.1_files/docker/registry/v2/repositories/prom/node-  
exporter/_manifests/tags/v0.16.0/index/sha256/55302581333c43d540db0e144cf9e7735423117  
a733cdec27716d87254221086/link  
1f3f02c4-fe58-4626-992b-  
c6c0fe4c8acf.34107.1_files/docker/registry/v2/repositories/prom/node-  
exporter/_layers/sha256/224a21997e8ca8514d42eb2ed98b19a7ee2537bce0b3a26b8dff510ab637f  
15c/link  
1f3f02c4-fe58-4626-992b-  
c6c0fe4c8acf.34107.1_files/docker/registry/v2/repositories/prom/node-  
exporter/_layers/sha256/528dda9cf23d0fad80347749d6d06229b9a19903e49b7177d5f4f58736538  
d4e/link  
1f3f02c4-fe58-4626-992b-  
c6c0fe4c8acf.34107.1_files/docker/registry/v2/repositories/prom/node-  
exporter/_layers/sha256/188af75e2de0203eac7c6e982feff45f9c340eaac4c7a0f59129712524fa2  
984/link
```

私有 registry 的运维操作

查询私有镜像中的 *images*

```
$ curl --user foo:foo123 --cacert /etc/docker/certs.d/172.27.132.67\:8000/ca.crt  
https://172.27.132.67:8000/v2/_catalog  
{  
  "repositories": ["prom/node-exporter"]  
}
```

查询某个镜像的 *tags* 列表

```
$ curl --user foo:foo123 --cacert /etc/docker/certs.d/172.27.132.67\:8000/ca.crt  
https://172.27.132.67:8000/v2/prom/node-exporter/tags/list  
{  
  "name": "prom/node-exporter",  
  "tags": ["v0.16.0"]  
}
```

获取 image 或 layer 的 digest

向 v2/<repoName>/manifests/<tagName> 发 GET 请求，从响应的头部 Docker-Content-Digest 获取 image digest，从响应的 body 的 fsLayers.blobSum 中获取 layDigests;

注意，必须包含请求头：Accept: application/vnd.docker.distribution.manifest.v2+json:

```
$ curl -v -H "Accept: application/vnd.docker.distribution.manifest.v2+json" --user
foo:foo123 --cacert /etc/docker/certs.d/172.27.132.67\8000/ca.crt
https://172.27.132.67:8000/v2/prom/node-exporter/manifests/v0.16.0
* About to connect() to 172.27.132.67 port 8000 (#0)
*   Trying 172.27.132.67...
* Connected to 172.27.132.67 (172.27.132.67) port 8000 (#0)
* Initializing NSS with certpath: sql:/etc/pki/nssdb
*   CAfile: /etc/docker/certs.d/172.27.132.67:8000/ca.crt
   Cpath: none
* SSL connection using TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256
* Server certificate:
*   subject: CN=registry,OU=4Paradigm,O=k8s,L=BeiJing,ST=BeiJing,C=CN
*   start date: Jul 05 12:52:00 2018 GMT
*   expire date: Jul 02 12:52:00 2028 GMT
*   common name: registry
*   issuer: CN=kubernetes,OU=4Paradigm,O=k8s,L=BeiJing,ST=BeiJing,C=CN
* Server auth using Basic with user 'foo'
> GET /v2/prom/node-exporter/manifests/v0.16.0 HTTP/1.1
> Authorization: Basic Zm9vOmZvbzEyMw==
> User-Agent: curl/7.29.0
> Host: 172.27.132.67:8000
> Accept: application/vnd.docker.distribution.manifest.v2+json
>
< HTTP/1.1 200 OK
< Content-Length: 949
< Content-Type: application/vnd.docker.distribution.manifest.v2+json
< Docker-Content-Digest:
sha256:55302581333c43d540db0e144cf9e7735423117a733cdec27716d87254221086
< Docker-Distribution-API-Version: registry/2.0
< Etag: "sha256:55302581333c43d540db0e144cf9e7735423117a733cdec27716d87254221086"
< X-Content-Type-Options: nosniff
< Date: Fri, 06 Jul 2018 06:18:41 GMT
<
{
  "schemaVersion": 2,
  "mediaType": "application/vnd.docker.distribution.manifest.v2+json",
  "config": {
    "mediaType": "application/vnd.docker.container.image.v1+json",
    "size": 3511,
    "digest":
"sha256:188af75e2de0203eac7c6e982feff45f9c340eaac4c7a0f59129712524fa2984"
  },
  "layers": [
    {
```



```

        "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
        "size": 2392417,
        "digest":
"sha256:224a21997e8ca8514d42eb2ed98b19a7ee2537bce0b3a26b8dff510ab637f15c"
    },
    {
        "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
        "size": 560703,
        "digest":
"sha256:cdb7590af5f064887f3d6008d46be65e929c74250d747813d85199e04fc70463"
    },
    {
        "mediaType": "application/vnd.docker.image.rootfs.diff.tar.gzip",
        "size": 5332460,
        "digest":
"sha256:528dda9cf23d0fad80347749d6d06229b9a19903e49b7177d5f4f58736538d4e"
    }
]

```

删除 *image*

向 `/v2/<name>/manifests/<reference>` 发送 DELETE 请求，reference 为上一步返回的 Docker-Content-Digest 字段内容：

```

$ curl -X DELETE --user foo:foo123 --cacert
/etc/docker/certs.d/172.27.132.67\8000/ca.crt
https://172.27.132.67:8000/v2/prom/node-
exporter/manifests/sha256:68effe31a4ae8312e47f54bec52d1fc925908009ce7e6f734e1b54a4169
081c5
$

```

删除 *layer*

向 `/v2/<name>/blobs/<digest>` 发送 DELETE 请求，其中 digest 是上一步返回的 `fsLayers.blobSum` 字段内容：

```
$ curl -X DELETE --user foo:foo123 --cacert
/etc/docker/certs.d/172.27.132.67\8000/ca.crt
https://172.27.132.67:8000/v2/prom/node-
exporter/blobs/sha256:a3ed95caeb02ffe68cdd9fd84406680ae93d633cb16422d00e8a7c22955b46d
4
$ curl -X DELETE --cacert /etc/docker/certs.d/172.27.132.67\8000/ca.crt
https://172.27.132.67:8000/v2/prom/node-
exporter/blobs/sha256:04176c8b224aa0eb9942af765f66dae866f436e75acef028fe44b8a98e04551
5
$
```

常见问题

login 失败 416

执行 <http://docs.ceph.com/docs/master/install/install-ceph-gateway/> 里面的 s3 test.py 程序失败:

```
[k8s@zhangjun-k8s-01 cert]$ python s3test.py
Traceback (most recent call last):
  File "s3test.py", line 12, in
    bucket = conn.create_bucket('my-new-bucket')
  File "/usr/lib/python2.7/site-packages/boto/s3/connection.py", line 625, in create_bucket
    response.status, response.reason, body)
boto.exception.S3ResponseError: S3ResponseError: 416 Requested Range Not Satisfiable
```

解决办法:

1. 在管理节点上修改 ceph.conf
2. ceph-deploy config push zhangjun-k8s-01 zhangjun-k8s-02 zhangjun-k8s-03
3. systemctl restart 'ceph-mds@zhangjun-k8s-03.service'
systemctl restart ceph-osd@0
systemctl restart 'ceph-mon@zhangjun-k8s-01.service'
systemctl restart 'ceph-mgr@zhangjun-k8s-01.service'

For anyone who is hitting this issue

set default pg_num and pgp_num to lower value(8 for example), or set mon_max_pg_per_osd to a high value in ceph.conf

radosgw-admin doesn't throw proper error when internal pool creation fails, hence the upper level error which is very confusing.

<https://tracker.ceph.com/issues/21497>

login 失败 503

```
[root@zhangjun-k8s-01 ~]# docker login 172.27.132.67:8000
```

```
Username: foo
```

```
Password:
```

```
Error response from daemon: login attempt to https://172.27.132.67:8000/v2/ failed with status: 503 Service Unavailable
```

原因： docker run 缺少 --privileged 参数；